

piecewise_testing

April 19, 2017

```
In [1]: import os
import pyopencl as pcl
import numpy as np
import scipy.stats as ss
import pandas as pd
import math
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

In [ ]: clfile = os.path.join('..', 'src', 'glm_gpu', 'cl', 'logistic.cl')
os.path.isfile(clfile)

In [ ]: os.environ['PYOPENCL_COMPILER_OUTPUT'] = '1'
device = pcl.get_platforms()[0].get_devices()[2]
cntx = pcl.Context([device])
queue = pcl.CommandQueue(cntx)

with open(clfile, 'r') as f:
    programs = pcl.Program(cntx, f.read()).build()

In [ ]: device.name

In [ ]: theta = np.array([0.5, 0.5], dtype=np.float32)

In [ ]: theta.dtype

In [2]: x0_1 = ss.norm(loc=10.0, scale=2.0)
x0_0 = ss.norm(loc=7.0, scale=2.0)
x1_1 = ss.norm(loc=5.0, scale=3.0)
x1_0 = ss.norm(loc=-5.0, scale=3.0)

nsamps=1024

X_1 = pd.DataFrame(index=range(nsamps),
                    columns=['x0', 'x1', 'y'])
X_0 = pd.DataFrame(index=range(nsamps),
                    columns=['x0', 'x1', 'y'])
```

```
X_1.loc[:, 'x0'] = x0_1.rvs(size=(nsamps,)).astype(np.float32)
X_1.loc[:, 'x1'] = x1_1.rvs(size=(nsamps,)).astype(np.float32)
X_1.loc[:, 'y'] = np.ones(shape=(nsamps,)).astype(np.float32)
```

```
X_0.loc[:, 'x0'] = x0_0.rvs(size=(nsamps,)).astype(np.float32)
X_0.loc[:, 'x1'] = x1_0.rvs(size=(nsamps,)).astype(np.float32)
X_0.loc[:, 'y'] = np.zeros(shape=(nsamps,)).astype(np.float32)
```

```
X_all = pd.concat((X_1, X_0), ignore_index=True)
```

```
X_all = X_all.reindex(np.random.permutation(X_all.index))
```

```
X = X_all.loc[:, ['x0', 'x1']]
y = X_all.loc[:, 'y']
```

```
In [ ]: X.values
```

```
In [ ]: cost_arr = np.zeros(shape=(X.shape[0],)).astype(np.float32)
temp_avg = np.zeros(shape=(4,)).astype(np.float32)
```

```
In [ ]: # compute cost
```

```
X_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, hostbuf=
y_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, hostbuf=
theta_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, host
cost_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_WRITE | pcl.mem_flags.COPY_HOST_PTR, host
temp_buf = pcl.Buffer(cntx, pcl.mem_flags.WRITE_ONLY, size=temp_avg.nbytes)
scratch_buf = pcl.LocalMemory(np.float32().nbytes * 256)
```

```
nrows = np.int32(X.shape[0])
ncols = np.int32(X.shape[1])
```

```
out = np.zeros(shape=(X.shape[0],), dtype=np.float32)
out_buf = pcl.Buffer(cntx, pcl.mem_flags.WRITE_ONLY, size=out.nbytes)
```

```
In [ ]:
```

```
In [ ]: sig_event = programs.sig(queue,
                                X.shape,
                                None,
                                X_buf,
                                theta_buf,
                                out_buf,
                                nrows,
                                ncols)
```

```
cost_event = programs.logistic_cost_ols(queue,
                                         X.shape,
                                         None,
```

```

X_buf,
theta_buf,
y_buf,
cost_buf,
nrows,
ncols)

# reduction_event = programs.matrix_row_mean(queue,
#                                           (1024,),
#                                           (256,),
#                                           cost_buf,
#                                           temp_buf,
#                                           scratch_buf,
#                                           nrows,
#                                           np.int32(1))

In [ ]: nrows, ncols, theta.shape

In [ ]: sig_event.wait()
        pcl.enqueue_copy(queue, out, out_buf).wait()

        cost_event.wait()
        pcl.enqueue_copy(queue, cost_arr, cost_buf).wait()

        # reduction_event.wait()
        # pcl.enqueue_copy(queue, temp_avg, temp_buf).wait()

In [50]: def sig(X, theta):
        lin = X.dot(theta)
        sig = 1.0 / (1.0 + np.exp(-lin))
        return sig

        def lr_cost(X, theta, y):
            est = sig(X, theta)
            log_est = np.log(est)
            cost = y*log_est + (1-y) * (1-log_est)
            cost *= -1
            return cost

        def cost_ss(est, actual):
            cost = ((est - actual)**2.0)
            cost /= 2.0
            return cost

In [ ]: my_out = sig(X.values, theta)
        my_cost = cost(my_out, y)

In [ ]: my_out.shape, out.shape

```

```

In [ ]: x = list(range(len(my_out)))
        plt.scatter(x, my_out)
        plt.scatter(x, out, color='r')
        plt.show()

In [ ]: x = list(range(len(my_cost)))
        plt.scatter(x, my_cost)
        plt.scatter(x, cost_arr)
        plt.show()

In [ ]: my_cost == cost_arr

In [ ]: X.values.dot(theta)

In [ ]: X.values

In [ ]: (7.93563604 + 5.08651495)/2.0

In [ ]: theta

In [ ]: out

In [ ]: my_out.reshape((my_out.size, )) == out

In [ ]: out == my_out

In [29]: # ok now let's check the gradient stuff
        def grad(X, theta, y):
            diff = y - sig(X, theta)
            g = np.zeros(X.shape)
            for i in range(X.shape[0]):
                for j in range(X.shape[1]):
                    g[i,j] = diff[i] * X[i,j]
            return g

In [ ]: y

In [ ]: theta = theta.reshape((theta.size,))
        grad(X.values, theta, y.values), sig(X.values, theta), y.values

In [ ]: gradient = np.zeros(shape=X.shape, dtype=np.float32)

In [ ]: X_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, hostbuf=
        y_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, hostbuf=
        theta_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR, host
        grad_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_WRITE | pcl.mem_flags.COPY_HOST_PTR , hos

In [ ]: grad_event = programs.logistic_gradient_ols(queue,
                                                    X.shape,
                                                    None,

```

```
X_buf,  
theta_buf,  
y_buf,  
grad_buf,  
nrows,  
ncols)
```

```
In [ ]: gradient
```

```
In [ ]: grad_event.wait()
```

```
In [ ]: pcl.enqueue_copy(queue, gradient, grad_buf).wait()
```

```
In [ ]: gradient
```

```
In [ ]: gradient == grad(X.values, theta, y.values)
```

```
In [ ]: gradient
```

```
In [ ]: grad(X.values, theta, y.values)
```

```
In [ ]: # let's double check I can get this to work
```

```
In [61]: def sig(X, theta):  
    lin = X.dot(theta)  
    sig = 1.0 / (1.0 + np.exp(-lin))  
    return sig  
  
def lr_cost(X, theta, y):  
    est = sig(X, theta)  
    log_est = np.log(est)  
    cost = y*log_est + (1-y) * (1-log_est)  
    cost *= -1  
    return cost  
  
def cost_ss(est, actual):  
    cost = ((est - actual)**2.0)  
    cost /= 2.0  
    return cost  
  
def grad(X, theta, y):  
    diff = y - sig(X, theta)  
    g = np.zeros(X.shape)  
    for i in range(X.shape[0]):  
        for j in range(X.shape[1]):  
            g[i,j] = diff[i] * X[i,j]  
    return g
```

```

def fit_params(X, y, theta):
    tol = 1e-5
    learning_rate = 1e-2
    costs = [np.inf]
    for i in range(0, 1000):
        my_cost = lr_cost(X, theta, y)
        my_cost = my_cost.sum()
        costs.append(my_cost)
        if abs(my_cost - costs[-2]) < tol:
            break
        else:
            my_gradient = grad(X, theta, y)
            my_gradient = my_gradient.mean(axis=0)
            theta = theta + learning_rate * my_gradient

    return theta, costs

```

```

In [66]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.6)
        new_t = np.random.normal(size=(X_train.shape[1],)).astype(np.float32)
        fitted_theta, outcost = fit_params(X_train.values, y_train.values, new_t)

```

```

In [67]: fitted_theta

```

```

Out[67]: array([ 0.03598584,  1.04095844])

```

```

In [68]: lr_cost(X_test.values, fitted_theta, y_test.values).sum()

```

```

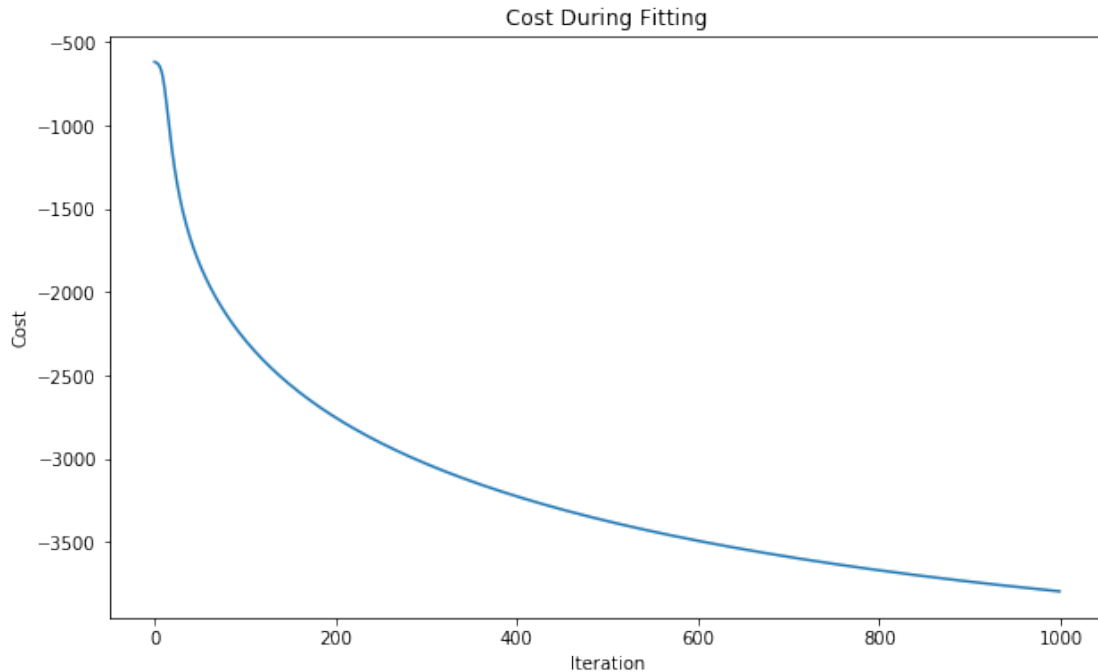
Out[68]: -2363.4378906858246

```

```

In [69]: plt.figure(figsize=(10,6))
        plt.plot(outcost[1:])
        plt.xlabel('Iteration')
        plt.ylabel('Cost')
        plt.title('Cost During Fitting')
        plt.show()

```



```
In [44]: ssc = cost_ss(sig(X_test.values, fitted_theta), y_test.values)
```

```
In [47]: ssc.sum()/X_test.shape[0]
```

```
Out[47]: 0.016629054794092435
```

```
In [ ]: cost(sig(X.values, fitted_theta), y).sum()
```

```
In [ ]: cost(sig(X.values, np.random.normal(size=theta.shape)), y).sum()
```

```
In [ ]: # ok that works, let's see if I can get this to work with the gpu code
```

```
def fit_gpu(X, y, theta, queue, cntx):
    tol = 1e-5
    learning_rate = 1e-2
    nrows = np.int32(X.shape[0])
    ncols = np.int32(X.shape[1])
    for i in range(0, 1000):
        cost_arr = np.zeros(shape=(X.shape[0],), dtype=np.float32)
        X_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR,
        y_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_PTR,
        theta_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_ONLY | pcl.mem_flags.COPY_HOST_P
        cost_buf = pcl.Buffer(cntx, pcl.mem_flags.READ_WRITE | pcl.mem_flags.COPY_HOST_P

        cost_event = programs.logistic_cost_ols(queue,X.shape, None,X_buf,theta_buf, y_bu
        cost_event.wait()
        pcl.enqueue_copy(queue, cost_arr, cost_buf)
```

```

cost_arr = cost_arr.sum()
cost_arr /= 2.
if cost_arr < tol:
    break
else:
    grad_buf = pcl.Buffer(cntx, pcl.mem_flags.WRITE_ONLY , size=X.nbytes)
    grad_event = programs.logistic_gradient_ols(queue,X.shape, None,X_buf,theta_b
    grad_event.wait()
    grad_arr = np.zeros(shape=X.shape, dtype=np.float32)
    pcl.enqueue_copy(queue, grad_arr, grad_buf)
    grad_arr = grad_arr.mean(axis=0)
    theta = theta + learning_rate * grad_arr

return theta

```

```

In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=0.6)
        new_t = np.random.normal(size=theta.shape).astype(np.float32)
        fitted_t_gpu = fit_gpu(X_train.values, y_train.values, new_t, queue, cntx)

```

```

In [ ]: fitted_t_gpu

```

```

In [ ]: fitted_theta = fit_params(X.values, y.values, new_t)

```

```

In [ ]: fitted_theta

```

```

In [ ]: fitted_t_gpu - fitted_theta

```

```

In [ ]:

```